

Jatin Shukla

AI/ML Engineer | Production ML | GenAI & Agentic Systems | MLOps

Gurugram, India | +91-9116237146 | shukla.jeetu2550@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

AI/ML engineer with ~3 years of experience building production ML systems, LLM applications, agentic workflows and data-driven AI products. Hands-on with Python, PyTorch, Scikit-learn, FastAPI, LangGraph/LangChain, SQL, BigQuery, Docker, CI/CD, Azure/GCP, model serving, monitoring and inference optimization. Experienced in taking AI solutions from data preparation and model prototyping through API integration, containerized deployment, UAT, monitoring and business adoption.

TECHNICAL SKILLS

AI/ML & GenAI: Python, PyTorch, TensorFlow, Scikit-learn, Hugging Face, LLMs, Generative AI, RAG, LangGraph, LangChain, LlamaIndex, Multi-Agent Systems, Tool Calling, Prompt Engineering, Embeddings

MLOps & Backend: FastAPI, Flask, REST APIs, ML Pipelines, Model Serving, Monitoring & Logging, CI/CD, Docker, Kubernetes Fundamentals, Git/GitHub, Joblib, Inference Optimization

Cloud & Data: Azure, GCP, BigQuery, AWS Fundamentals, SQL, ETL/ELT, Data Pipelines, Data Validation, Feature Engineering, FAISS, BM25, Cross-Encoder Reranking

EXPERIENCE

Senior Executive - AI

Sep 2026 - Present

JB Emporium India | Gurugram, India

- Build AI and data prototypes for enterprise sourcing and product workflows, using Python, SQL, BigQuery and LLM-based tools to turn messy multi-source data into structured, searchable and traceable outputs.
- Develop agentic and retrieval-driven workflows with validation, transparent filters and source-linked results; work with business stakeholders to convert ambiguous requirements into testable, production-oriented AI solutions.

Associate Research Engineer

Jan 2025 - Aug 2026

Havells India Ltd. | India

- Developed and deployed PyTorch/Scikit-learn models for fan air-delivery, velocity and pressure prediction, owning data preparation, feature engineering, model evaluation and repeatable inference workflows; improved engineering prediction accuracy by 18-22%.
- Built LED thermal-prediction pipelines from preprocessing and validation through model inference, reducing design-evaluation turnaround by 35% and making model outputs easier for engineering teams to consume.
- Optimized a Tri-Branch PointNet production inference pipeline using standardized preprocessing, feature scaling and Joblib model packaging, reducing end-to-end latency by 35% while keeping inference reproducible across design inputs.
- Productionized AI/ML capabilities behind FastAPI/Flask REST services and containerized deployment workflows, integrating model outputs into internal applications used by 30+ engineers; supported Git-based CI/CD and Azure deployment practices.
- Delivered 4+ LLM/GenAI and agentic solutions across HR, market intelligence and productivity workflows using retrieval, memory, tool routing, structured scoring, logging/traceability and human review, reducing manual research and screening effort by 40%+.
- Worked across product, engineering and business teams to move prototypes toward production: defined validation checks, UAT and fallback flows, documented model/API behavior, investigated failures and iterated on performance, reliability and user feedback.

SELECTED PROJECTS

AgentFlow - Multi-Agent AI Platform | [GitHub](#) - FastAPI, LangGraph, LangChain, Docker

- Built a 6-stage orchestration workflow with supervisor planning, specialist agents, memory retrieval, tool execution, reviewer scoring, trace persistence and human approval.
- Added guarded tool loops, structured run history and JSON trace export to make agent execution easier to inspect, debug and audit across repeated workflows.

Nexora - RAG Knowledge Assistant | [GitHub](#) - FastAPI, FAISS, BM25, Cross-Encoder

- Built a grounded RAG pipeline for PDFs and web content covering ingestion, chunking, embeddings, hybrid retrieval, Cross-Encoder reranking and source-backed generation.
- Implemented persistent source/chunk storage, multi-mode Q&A and retrieval-focused evaluation so responses stay traceable to the underlying knowledge base.

Agentic AI Code Review Bot | [GitHub](#) - FastAPI, Docker, GitHub API, Multi-Agent

- Built an 8-agent PR review system covering bugs, security, code quality and test coverage with severity scoring, risk summaries and human-reviewable fixes.
- Integrated GitHub APIs with a containerized FastAPI backend and structured review outputs, keeping merge decisions human-controlled while reducing repetitive PR triage.

EDUCATION

M.Tech CSE - NIT Hamirpur | CGPA 9.15/10

2023 - 2025

B.Tech - IET Lucknow | CGPA 7.8/10

2018 - 2022